



E9: 309 Advanced Deep Learning

19-10-2020

Instructor: Sriram Ganapathy
sriramg@iisc.ac.in

Teaching Assistant : Akshara Soman, Prachi Singh, Jaswanth Reddy
aksharas@iisc.ac.in, prachis@iisc.ac.in, jaswanthk@iisc.ac.in

<http://leap.ee.iisc.ac.in/sriram/teaching/ADL2020/>

Housekeeping

✳ Filling the google form in the webpage

➡ Contents will be made available to the folks in the creditors mailing lists.

✓ Announcements regarding evaluations and projects will be shared only with creditors as well as video links.

★ Teams channel interaction and TA session for creditors only.

✳ Online registration portal from academics.iisc.ac.in

✓ Your research/faculty advisor may need to approve also before the deadline (Oct. 20th?)



Recap of previous class



Some notations

✳ $\mathbf{x} \in \mathcal{R}^D$ - input data.


✳ $\mathbf{y} \in \mathcal{R}^C$ - neural network targets.

✳ $\hat{\mathbf{y}} \in \mathcal{B}^C$ - model outputs.

✳ $\mathbf{e}, \mathbf{h} \in \mathcal{R}^d$ - hidden model representations or embeddings.

✳ Θ - collection of learnable parameters in the model.

✳ $E(\mathbf{y}, \hat{\mathbf{y}})$ - error function used in the model training.



Some notations

✳ $\{\mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{y}_1, \dots, \mathbf{y}_N\}$ - labeled training data

✳ $q = \{1 \dots Q\}$ - iteration index.

✳ $t = \{1 \dots T\}$ - discrete time index.


✳ $l = \{1 \dots L\}$ - layer index

✳ η - learning rate (hyper-parameter)


✳ N_b - mini-batch size and B is the number of mini-batches.



Long-term dependency issues



Long short term memory (LSTM) idea



State of affairs

✳ Recurrent networks

- ➡ First order recurrence

- ➡ Back propagation in RNNs

✳ Issues with long-term dependency

- ➡ Gates as neural networks

- ➡ LSTM and GRUs


Reading Assignment - “A review of recurrent neural networks - LSTM cells and Network Architectures”

https://www.mitpressjournals.org/doi/pdfplus/10.1162/neco_a_01199




Other network architectures

* Multiple input multiple output




Bidirectional RNNs

🌟 Multiple input multiple output



Training data



$$X = \{x^{(1)} \dots x^{(T)}\} \longrightarrow Y$$


Bidirectional RNNs

- * Forward and backward recurrence variables are summed *or concatenated*
- * Can we implemented using LSTMS - called BLSTMs
- * Backpropagation
 - Gradients for the forward recurrence will have a backward relationship
 - Gradients for the backward recurrence will have a forward relationship
- * Commonly used in offline processing of sequence data
 - Mostly improves over the forward LSTM/RNN models.



Other network architectures


✳ Multiple input single output (seq2vec)



✳ Applications like topic summarization of text, speaker identification of speech.

Other network architectures

* Single input multiple output



* Image captioning.

Encoder-decoder models

✳ Multiple input multiple output (with different label index) - Seq2seq

$$\begin{array}{c} x(1) \dots x(T) \\ \downarrow \\ \underline{y(1)} \dots \underline{y(S)} \end{array}$$

$$\underline{x} = \{x_1 \dots x_D\}$$

$x(1)$ ✗

✳ Applications

✓ Machine translation.


✓ Speech recognition.

✓ Video captioning.




Encoder-decoder models

✳ Multiple input multiple output (with different label index) - Seq2seq



Encoder-decoder models

✳ Multiple input multiple output (with different label index) - Seq2seq



Encoder-decoder models

* Encoder — convert sequence $\mathbf{x} = \{\mathbf{x}(1), \dots, \mathbf{x}(T)\}$ to vector

$$\mathbf{h}(t) = f(\mathbf{h}(t-1), \mathbf{x}(t))$$

$$\mathbf{e} = f'(\mathbf{h}_1, \dots, \mathbf{h}_T)$$

* The encoder can have multiple deep RNN layers.

* For simplicity

$$\mathbf{e} = \mathbf{h}_T$$



Encoder-decoder models

* Encoder — convert sequence $\mathbf{x} = \{\mathbf{x}(1), \dots, \mathbf{x}(T)\}$ to vector $\mathbf{e}$

* Decoder — converts the vector embedding from the encoder to the output sequence $\hat{\mathbf{y}} = \{\hat{\mathbf{y}}(1), \dots, \hat{\mathbf{y}}(S)\}$ with different label index.

$$p(\hat{\mathbf{y}}) = \prod_{s=1}^S p(\hat{\mathbf{y}}(s) | \hat{\mathbf{y}}(1), \dots, \hat{\mathbf{y}}(s-1))$$

$$\hat{\mathbf{y}} = \hat{y} = \{\hat{y}(1), \dots, \hat{y}(s)\}$$

* RNN decoder assumption

$$p(\hat{\mathbf{y}}(s) | \hat{\mathbf{y}}(1), \dots, \hat{\mathbf{y}}(s-1)) = p(\hat{\mathbf{y}}(s) | \hat{\mathbf{y}}(s-1), \mathbf{e}) = \text{softmax}(\mathbf{V}\hat{\mathbf{y}}(s-1) + \mathbf{R}\mathbf{c}(s-1) + \mathbf{T}\mathbf{e} + \mathbf{d})$$
$$\mathbf{c}(s) = f(\mathbf{c}(s-1), \mathbf{e})$$

* The decoder can also have multiple layers of deep RNNs before softmax.



Encoder-decoder models

- ✳ Encoder — convert sequences to vectors

- ✳ Decoder — converts the vector embedding from the encoder to the output sequence with different label index.

 - ✓ Start and end label are also encoded as output vector indices.

 - ★ Enable the starting and ending of the output sequence.

- ✳ Assumption

 - ✓ The entire input sequence can be represented as a single vector e

 - ★ May not be able to perform this efficiently for long sequences.



Encoder-decoder models

✳️ Modification of encoder-decoder model

$$p(\hat{\mathbf{y}}(s)|\hat{\mathbf{y}}(1), \dots, \hat{\mathbf{y}}(s-1)) = p(\hat{\mathbf{y}}(s)|\hat{\mathbf{y}}(s-1), \mathbf{e}) = \text{softmax}(\mathbf{V}\hat{\mathbf{y}}(s-1) + \mathbf{R}\mathbf{c}(s-1) + \mathbf{T}\mathbf{e} + \mathbf{d})$$

$$p(\hat{\mathbf{y}}(s)|\hat{\mathbf{y}}(1), \dots, \hat{\mathbf{y}}(s-1)) = p(\hat{\mathbf{y}}(s)|\hat{\mathbf{y}}(s-1), \mathbf{e}(s)) = \text{softmax}(\mathbf{V}\hat{\mathbf{y}}(s-1) + \mathbf{R}\mathbf{c}(s-1) + \mathbf{T}\mathbf{e}(s) + \mathbf{d})$$

✳️ where

$$\mathbf{e}(s) = \sum_{t=1}^T \alpha(s, t) \mathbf{h}(t)$$

✳️ Here $\alpha(s, t)$ captures the contribution of input at time t with output at time s




Encoder-decoder models

✳ Obtaining the relative contribution $\alpha(s, t)$

✓ Implementing this automatically using network-in-network

Attention network



$$\hat{\alpha}(s, t) = \mathbf{A}[\mathbf{c}(s-1); \mathbf{h}(t)]$$

$$\alpha(s, t) = S(\hat{\alpha}(s, t)) = \frac{\exp(\hat{\alpha}(s, t))}{\sum_{t'} \exp(\hat{\alpha}(s, t'))}$$

✓ The values $\alpha(s, t)$ are called attention weights.


Analysis of attention networks

✱ Attention weights $\alpha(s, t)$

- ✓ Probability of linking (attending) to input at t for generating output at s
- ✓ Useful in analyzing the internal structure of the encoder-decoder model

Visualizing the attention weights

Reading Assignment - “Neural Machine Translation by Jointly Learning to Align and Translate”
<https://arxiv.org/pdf/1409.0473.pdf>




Analysis of attention networks

✳ Attention weights $\alpha(s, t)$

- ✓ Probability of linking (attending) to input at t for generating output at s
- ✓ Useful in analyzing the internal structure of the encoder-decoder model

Visualizing the attention weights

Reading Assignment - “Neural Machine Translation by Jointly Learning to Align and Translate”
<https://arxiv.org/pdf/1409.0473.pdf>



Visualizing attention

