

E9: 309 ADL 18-11-2020



Housekeeping

🌟 1st mini-project

✓ Deadlines

★ Presentation on Nov19 and Nov20

★ Your date allocation has been finalized

Teams assignment link.

★ Presentation and report template will be sent out this week.

★ Report 1 page + references and tools

★ Slides 4 slides for individual project and 6 slides for 2-member.

Attachments (Deadline : Nov 19: 10AM)
Firstname - Lastname - proj1-slides.
ppt

Firstname - Lastname - proj1-report
.pdf



Recap of previous class



Representation learning/data-visualization

* Restricted Boltzmann machine

→ Energy based model

✓ Conditional independence

✓ Sigmoidal function for conditional probability

$$p(h_j = 1/v) = \sigma(\cdot)$$

→ Issues in RBM training



RBM - Training

* Model parameters $\Theta = \{\mathbf{W}, \mathbf{b}, \mathbf{c}\}$ $\mathbf{x} = [\mathbf{v}^T \mathbf{h}^T]^T$

✓ Learnt by maximizing the log-likelihood $\log(p(\mathbf{x}; \Theta))$

✓ Non-convex optimization.

* Gradient ~~descent~~ ^{as cnt} based optimization

$$\frac{\partial \log(p(\mathbf{x}; \Theta))}{\partial \Theta} = \frac{\partial \log(\tilde{p}(\mathbf{x}; \Theta))}{\partial \Theta} - \frac{\partial \log(Z(\Theta))}{\partial \Theta}$$

$$\frac{\partial \log(Z(\Theta))}{\partial \Theta} = \frac{1}{Z} \frac{\partial Z(\Theta)}{\partial \Theta}$$

$$Z(\Theta) = \sum_{\mathbf{x}} \tilde{p}(\mathbf{x}; \Theta)$$



RBM - Training

✳ For exponential families

$$\tilde{p}(\mathbf{x}; \Theta) > 0 \quad \forall \mathbf{x}$$

$$\frac{\partial \log(Z(\Theta))}{\partial \Theta} = \frac{1}{Z} \sum_{\mathbf{x}} \frac{\partial \tilde{p}(\mathbf{x}; \Theta)}{\partial \Theta}$$

$$\frac{\partial \log(Z(\Theta))}{\partial \Theta} = \frac{1}{Z} \sum_{\mathbf{x}} \frac{\partial \exp(\log(\tilde{p}(\mathbf{x}; \Theta)))}{\partial \Theta}$$

$$= \frac{1}{Z} \sum_{\mathbf{x}} \exp(\log(\tilde{p}(\mathbf{x}; \Theta))) \frac{\partial (\log(\tilde{p}(\mathbf{x}; \Theta)))}{\partial \Theta}$$

$$= \sum_{\mathbf{x}} p(\mathbf{x}; \Theta) \frac{\partial (\log(\tilde{p}(\mathbf{x}; \Theta)))}{\partial \Theta}$$



RBM - Training

$$\frac{\partial(\log Z(\Theta))}{\partial \Theta} = \mathbb{E}_{\mathbf{x} \sim p} \left[\frac{\partial \tilde{p}(\mathbf{x}, \Theta)}{\partial \Theta} \right]$$

→ Intractable

✱ Intractable to compute the exact gradient of the negative phase

➡ Using approximations to gradients

✓ Based on sampling methods.

★ Monte-carlo Markov Chain (MCMC) based approximation

★ Resorting to Gibbs sampling.



Approximating expectations

* Expectation is intractable in the gradient computation.

$$s = \mathbb{E}_{x \sim p} [f(\mathbf{x})]$$
$$\hat{s}_n = \frac{1}{n} \sum_{i=1}^n f(\mathbf{x}_i)$$

$f(\mathbf{x}) = \frac{\partial}{\partial \theta} \log \tilde{p}(\mathbf{x}; \theta)$

$\hat{s}_n \rightarrow s$

* Using law of large numbers and central limit theorem

$$\mathbb{E}[\hat{s}_n] = s ; \quad \lim_{n \rightarrow \infty} \hat{s}_n \Rightarrow s ; \quad \underbrace{\text{Var}[\hat{s}_n]} = \underbrace{\frac{\text{Var}[f(\mathbf{x})]}{n}}$$

$x_i?$

every step of gradient ascent

$$\theta = \{w, b, c\}$$

↑ gradient ascent


gradient
negative
phase

← approx.

$$p(x; \theta)$$

↓ sampling
 n

$$\{\tilde{x}_i\}_{i=1}^n$$



Sampling

* The question of finding the right samples $\mathbf{x}_i$

* Should we sample based on the p or some other function q

$$p(\mathbf{x})f(\mathbf{x}) = q(\mathbf{x}) \frac{p(\mathbf{x})f(\mathbf{x})}{q(\mathbf{x})}$$

* Using a suitable function q , the estimate of the intractable expectation is

$\hat{s}_q$

$$\hat{s}_q = \frac{1}{n} \sum_{i=1, \mathbf{x}_i \sim q}^n \frac{p(\mathbf{x}_i)f(\mathbf{x}_i)}{q(\mathbf{x}_i)}$$

← $p(\mathbf{x}_i)$ ← $f(\mathbf{x}_i)$
→ $q(\mathbf{x}_i)$

Infeasible RBMs

Importance Sampling



Gibbs sampling with random initialization

✳ Using Markov chain Monte-Carlo (MCMC) sampling $p_{model}(\mathbf{x})$

- ✓ Initialize random samples using uniform distribution.
- ✓ Use the data samples to perform new updates of the samples
- ✓ Perform k steps of this update.

$$\tilde{\mathbf{x}} \leftarrow \underline{\underline{p(\mathbf{x}; \theta)}}$$

$$\underline{\underline{\tilde{\mathbf{x}}^0}} \rightarrow \tilde{\mathbf{x}}^1 \rightarrow \tilde{\mathbf{x}}^2 \rightarrow \dots$$

k-steps

Markov approx.



Gibbs sampling with random initialization

Algorithm 18.1 A naive MCMC algorithm for maximizing the log-likelihood with an intractable partition function using gradient ascent.

Set ϵ , the step size, to a small positive number.

Set k , the number of Gibbs steps, high enough to allow burn in. Perhaps 100 to train an RBM on a small image patch.

while not converged **do**

Sample a minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from the training set.

$\mathbf{g} \leftarrow \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log \tilde{p}(\mathbf{x}^{(i)}; \theta)$ *positive phase*

Initialize a set of m samples $\{\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(m)}\}$ to random values (e.g., from a uniform or normal distribution, or possibly a distribution with marginals matched to the model's marginals).

for $i = 1$ to k **do**

for $j = 1$ to m **do**

$(\tilde{\mathbf{x}}^{(j)}) \leftarrow \text{gibbs_update}(\tilde{\mathbf{x}}^{(j)})$.

end for

end for

$\mathbf{g} \leftarrow \mathbf{g} - \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log \tilde{p}(\tilde{\mathbf{x}}^{(i)}; \theta)$ *negative phase*

$\theta \leftarrow \theta + \epsilon \mathbf{g}$.

end while

Chap 18, 20
of
Deep Learning
Book

Gibbs sampling



Gibbs sampling with random initialization

✱ Using Markov chain Monte-Carlo (MCMC) sampling $p_{model}(\mathbf{x})$

✓ Use the given data samples and the current model weights to perform Gibbs sampling

★ Initially the model weights make for a poor estimation of the negative phase of the gradient.

★ But the positive phase makes up for the lossy estimate.

★ Once the model weights are updated for a few iterations

★ The negative phase becomes more accurate.

$x_1 \dots x_N$
↓
neighborhood
 $\tilde{x}_1 \dots \tilde{x}_N$

Contrastive Divergence



Gibbs sampling with random initialization

Algorithm 18.2 The contrastive divergence algorithm, using gradient ascent as the optimization procedure.

Set ϵ , the step size, to a small positive number.

Set k , the number of Gibbs steps, high enough to allow a Markov chain sampling from $p(\mathbf{x}; \theta)$ to mix when initialized from p_{data} . Perhaps 1-20 to train an RBM on a small image patch.

while not converged **do**

 Sample a minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from the training set.

$\mathbf{g} \leftarrow \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log \tilde{p}(\mathbf{x}^{(i)}; \theta)$.

for $i = 1$ to m **do**

$\tilde{\mathbf{x}}^{(i)} \leftarrow \mathbf{x}^{(i)}$.

end for

for $i = 1$ to k **do**

for $j = 1$ to m **do**

$\tilde{\mathbf{x}}^{(j)} \leftarrow \text{gibbs_update}(\tilde{\mathbf{x}}^{(j)})$.

end for

end for

$\mathbf{g} \leftarrow \mathbf{g} - \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log \tilde{p}(\tilde{\mathbf{x}}^{(i)}; \theta)$.

$\theta \leftarrow \theta + \epsilon \mathbf{g}$.

end while

+ phase

Contrastive
divergence
k-step.

-ve phase



One step contrastive divergence - Training example

➔ Given a set of visible data $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_N\} \in \mathcal{B}^D$

$$\mathbf{x} = [\mathbf{v}^T \ \mathbf{h}^T]^T$$

➔ Randomly initialize the model parameters $\Theta^0 = \{\mathbf{W}^0, \mathbf{b}^0, \mathbf{c}^0\}; \underline{k=0}$ ^{iteration}

✓ Sampling $\{\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_N\} \in \mathcal{B}^d$

◦ Using conditional independence of hidden given visible

★ Sample each $\{\tilde{\mathbf{h}}_1, \tilde{\mathbf{h}}_2, \dots, \tilde{\mathbf{h}}_N\} \in \mathcal{B}^d$

$$p(\underline{h_{q,j}} = 1 | \underline{\mathbf{v}_q}) = \sigma(\underline{\mathbf{v}_q^T \mathbf{W}_{:,j}^k} + \underline{c_j^k}) \quad \xrightarrow{\sim h_q} \quad q = \{1, \dots, N\}, j = \{1, \dots, d\}$$



One step contrastive divergence - Training example

➔ Given a set of visible data $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_N\} \in \mathcal{B}^D$ $\mathbf{x} = [\mathbf{v}^T \ \mathbf{h}^T]^T$

➔ Randomly initialize the model parameters $\Theta^0 = \{\mathbf{W}^0, \mathbf{b}^0, \mathbf{c}^0\}$ $k = 0$

✓ Sampling visible nodes again using $[\tilde{\mathbf{h}}_1, \tilde{\mathbf{h}}_2, \dots, \tilde{\mathbf{h}}_N]$

◦ Using conditional independence of visible given hidden

★ Sample visible nodes $\{\tilde{\mathbf{v}}_1, \tilde{\mathbf{v}}_2, \dots, \tilde{\mathbf{v}}_N\} \in \mathcal{B}^D$

$$p(v_{q,i} = 1 | \tilde{\mathbf{h}}_q) = \sigma(\mathbf{W}_{i,:}^k \tilde{\mathbf{h}}_q^T + b_i^k) \quad q = \{1, \dots, N\}, i = \{1, \dots, D\}$$

$$\begin{bmatrix} \tilde{\mathbf{v}}_1 & \tilde{\mathbf{v}}_2 & \dots & \tilde{\mathbf{v}}_N \\ \tilde{\mathbf{h}}_1 & \tilde{\mathbf{h}}_2 & \dots & \tilde{\mathbf{h}}_N \end{bmatrix}$$

$\tilde{\mathbf{v}}_1 \quad \tilde{\mathbf{v}}_2 \quad \dots \quad \tilde{\mathbf{v}}_N$



One-step contrastive divergence

* Computing the gradient

observed data

positive phase

$$\frac{\partial(p([\mathbf{v} \ \mathbf{h}], \Theta)}{\partial \mathbf{W}} \approx \frac{1}{N} \sum_{q=1}^N \mathbf{v}_q \mathbf{h}_q^T - \frac{1}{N} \sum_{q=1}^N \tilde{\mathbf{v}}_q \tilde{\mathbf{h}}_q^T$$

model generated data

model



Handwritten notes:

- $\log \tilde{p}(\mathbf{x}; \Theta)$
- $-E[\mathbf{v}, \mathbf{h}]$
- $\cdot \mathbf{v}^T \mathbf{W} \mathbf{h}$
- $\log \tilde{p}(\tilde{\mathbf{x}}; \Theta)$

* Performing gradient ascent using the approximate gradient

$$\Theta^{k+1} = \Theta^k + \eta \left. \frac{\partial \log(p(\mathbf{x}, \Theta))}{\partial \Theta} \right|_{\Theta = \Theta^k}$$

Positive phase and negative phase



Convergence $\approx$ +ve phase is the same -ve samples

"Motivation"

✳️ Neuroscientific motivation of the learning in the brain

➡️ Real world samples in our day-to-day [Positive phase] ↑

➡️ Hallucinations and dreams [Negative phase] suppress ↓

✳️ Learning from the data.

➡️ Balancing the positive phase based learning with the negative phase.

Gaussian Bernoulli RBM

v - continuous

h - discrete

✳ For modeling real observations $\mathbf{v} \in \mathcal{R}^D$

✳ Define the energy function

$$E[\mathbf{v}, \mathbf{h}] = \frac{1}{2}(\mathbf{v} - \mathbf{a})^T(\mathbf{v} - \mathbf{a}) - \mathbf{v}^T \mathbf{W} \mathbf{h} - \mathbf{b}^T \mathbf{h}$$

$$p([\mathbf{v}, \mathbf{h}]) = \frac{e^{-E[\mathbf{v}, \mathbf{h}]}}{Z}$$

✳ The conditional distributions

$$p(\mathbf{v}|\mathbf{h}) = \mathcal{N}(\mathbf{W} \mathbf{h} + \mathbf{a}, \mathbf{I})$$

$$p(h_j = 1|\mathbf{v}) = \sigma(\mathbf{v}^T \mathbf{W}_{:,j}^k + b_j)$$

Take home exercise.



Properties of GRBM

✳ $d = 0$


Marginal

$$E(\mathbf{v}) = \frac{1}{2}(\mathbf{v} - \mathbf{a})^T(\mathbf{v} - \mathbf{a})$$

$$p(\mathbf{v}) = \frac{e^{-E(\mathbf{v})}}{Z}$$

✳ The marginal distribution is a Gaussian. =

$$N(\underline{a}, \underline{I})$$



Properties of GRBM

* $d = 1$

$$E[\mathbf{v}, h] = \frac{1}{2}(\mathbf{v} - \mathbf{a})^T(\mathbf{v} - \mathbf{a}) - \underbrace{h\mathbf{v}^T\mathbf{w}} - \underbrace{hb}$$

$$p([\mathbf{v}, h]) = \frac{e^{-E[\mathbf{v}, h]}}{Z}$$

$$\rightarrow p([\mathbf{v}, h = 0]) = \underline{\alpha \mathcal{N}(\mathbf{a}, \mathbf{I})}$$


$$p([\mathbf{v}, h = 1]) = (1 - \alpha) \mathcal{N}(\underline{\mathbf{a} + \mathbf{w}}, \mathbf{I})$$

* The marginal distribution is then

$$\alpha \mathcal{N}(\mathbf{a}, \mathbf{I}) + (1 - \alpha) \mathcal{N}(\underline{\mathbf{a} + \mathbf{w}}, \mathbf{I})$$

$$\underline{p(\mathbf{v}) = p([\mathbf{v}, h = 0]) + p([\mathbf{v}, h = 1])}$$

✓ 2-mixture Gaussian





$h = 0$
 $h = 1$

Properties of GRBM

* For any general d dimensions

$E_{\mathbb{R}^D}$



$$\rightarrow E[\mathbf{v}, \underline{\mathbf{h}}_d] = E[\mathbf{v}, \underline{\mathbf{h}}_{d-1}] + \underline{h_d} \mathbf{v}^T \underline{\mathbf{W}}_{:,d} + \underline{b_d h_d}$$

$$p([\mathbf{v}, [\underline{\mathbf{h}}_{d-1}, h_d = 0]]) = \alpha p([\mathbf{v}, \underline{\mathbf{h}}_{d-1}])$$

$$p([\mathbf{v}, [\underline{\mathbf{h}}_{d-1}, h_d = 1]]) = (1 - \alpha) p([\underline{\mathbf{v}} + \underline{\mathbf{W}}_{:,d}, \underline{\mathbf{h}}_{d-1}])$$

$\underline{D \times d}$
 $\underline{h_d = \begin{Bmatrix} 0 \\ 1 \end{Bmatrix}}$

* For d=0, 1 Gaussian, d=1, 2-mix Gaussian, ...

→ 2^d mixture Gaussian for any arbitrary d.


$$\mathbf{h}_d = \begin{bmatrix} \mathbf{h}_{d-1} \\ h_d \end{bmatrix}$$



GRBMs and GMMs


$p(v)$

(a)




$d=0$

(b)




$d=1$

(c)



(d)



$d=3$

8 - Gaussian




Deep Belief Networks (DBN)


✳ Stacking RBMs in a disjoint fashion

- ✓ Layer-wise training for each RBM.
 - Weights are frozen each layer before training the next layer.
- ✓ Ancestral sampling can be performed for data generation
 - ★ Lossy sample generation due to accumulation of errors.


✳ Most common use - pre-training of DNNs.




DBNs for initialization



DBNs for visualization



PCA



RBM



More reading

Salakhutdinov, Ruslan, Andriy Mnih, and Geoffrey Hinton. "Restricted Boltzmann machines for collaborative filtering." *Proceedings of the 24th international conference on Machine learning*. 2007.





Data generation using RBMs

learning



generating



Data generation using RBMs



Source — <https://lme.tf.fau.de/lecture-notes/lecture-notes-in-deep-learning-unsupervised-learning-part-1/>

Deep Boltzmann machine

* Deep layers of connections with RBM structure.

→ Joint energy function.

→ Undirected graph

