

Deep Learning - Theory and Practice

Linear Regression, Least Squares
Classification and Logistic Regression

20-02-2020

<http://leap.ee.iisc.ac.in/sriram/teaching/DL20/>

deeplearning.cce2020@gmail.com



Least Squares for Classification

- ❖ K-class classification problem

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$$

- ❖ With 1-of-K hot encoding, and least squares regression

$$\mathbf{t} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

$$\sum_{n=1}^N$$


$$E_D(\tilde{\mathbf{W}}) = \frac{1}{2} \text{Tr} \left\{ (\tilde{\mathbf{X}} \tilde{\mathbf{W}} - \mathbf{T})^T (\tilde{\mathbf{X}} \tilde{\mathbf{W}} - \mathbf{T}) \right\}$$

$$N \times N$$

Bishop - PRML book (Chap 3)



$$T_H \begin{pmatrix} \|y(x_1) - t_1\|^2 \\ \|y(x_2) - t_2\|^2 \\ \vdots \\ \|y(x_N) - t_N\|^2 \end{pmatrix}_{N \times N}$$



Logistic Regression (classifier)

- ❖ 2- class logistic regression

$$p(C_1|\phi) = y(\phi) = \sigma(\mathbf{w}^T \phi)$$

- ❖ Maximum likelihood solution

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \phi_n$$

- ❖ K-class logistic regression

$$p(C_k|\phi) = y_k(\phi) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$$

- ❖ Maximum likelihood solution

$$a_k = \mathbf{w}_k^T \phi.$$

$$\nabla_{\mathbf{w}_j} E(\mathbf{w}_1, \dots, \mathbf{w}_K) = \sum_{n=1}^N (y_{nj} - t_{nj}) \phi_n$$

Issues

- * Outputs are real numbers
- * Model targets are binary.
- * Cost Function - least squares

Next Model. → make outputs
close to prob.

$$P(C_1/\phi(x)) = \frac{\text{Binary Classification } (C_1, C_2)}{P(\phi(x)/C_1) P(C_1)}$$

$$\phi(x) \quad P(\phi(x)) = P(\phi(x), C_1) + P(\phi(x), C_2)$$

↑ fixed non-linear transformation

$$P(C_1/\phi(x)) = \frac{1}{1 + e^{-f(x)}} ; \quad \frac{1}{1 + \frac{P(\phi(x), C_2)}{P(\phi(x), C_1)}}$$


$$\boxed{\underline{\underline{f(x)}} = -\ln \frac{P(\phi(x), C_2)}{P(\phi(x), C_1)}}$$

$$p(C_1/\phi(x)) = \frac{1}{1 + e^{-f(x)}} = \underset{\substack{\uparrow \\ \text{logistic} \\ \text{function}}}{\sigma}(f(x))$$


Approx.

$$f(x) = \underline{\underline{\underline{\underline{w}}}}^T \phi(x) + w_0 \quad \underline{\underline{\underline{\tilde{w}}}}} = [\underline{\underline{\underline{w}}}^T \ w_0]^T$$

Model



$\sigma(a)$



$$p(C_2 | \phi(x)) = 1 - \sigma(f(x))$$

$$= \sigma(-f(x))$$

$$t_n = 0 \text{ if } x_n \in C_2$$

$$= 1 \text{ if } x_n \in C_1$$

$$y(x) = \underline{\sigma(f(x))} = p(C_1 | \phi(x)) \quad x_1, x_2 \dots x_N.$$

Likelihood

$$\underset{\substack{\uparrow \\ \omega}}{L(X)} = \prod_{n=1}^N (y(\underline{x_n}))^{t_n} (1 - y(x_n))^{1-t_n}$$

w.r.t $\underline{w}, w_0$

$$\underline{\log L(x)} = \sum_{n=1}^N t_n \log y_n + (1-t_n) \log (1-y_n)$$

B.C.E (binary cross entropy) = $-\log L(x)$
 between $\begin{pmatrix} y(x_n) \\ 1-y(x_n) \end{pmatrix}, \begin{pmatrix} t_n \\ 1-t_n \end{pmatrix}$

$y_n = y(x_n)$

$$y_n = y(x_n) = \frac{1}{1 + e^{-\{ \underline{w}^T \phi(x_n) + \underline{w_0} \}}}$$


Minimize B.C.E
 w.r.t w, w_0

$$X = \{x_1 \dots x_N\}$$

Typical Error Surfaces


B.C.E
↑

Typical Error Surface as a function of parameters (weights and biases)




E_{total}


w_0



A hand-drawn diagram of a non-convex error surface, showing a wavy line with multiple peaks and valleys. The label E_{total} is written above the line, and w_0 is written below it. An arrow points from the label w_0 to the right.



Non-convex
Multiple max and min




zero saddle point

Iterative

$$E(w^{t+1}) \leq E(w^t)$$

$$\frac{\partial E}{\partial w}$$



Direction (Error surfaces are smooth)

Derivative $\frac{\partial E}{\partial w}$ (-ve direction of derivative)



Magnitude - η (learning rate) +ve

$$w^{t+1} = w^t - \eta \left. \frac{\partial E}{\partial w} \right|_{w=w^t}$$



Gradient Descent Algorithm

Learning with Gradient Descent

Error surface close to a local




Learning Using Gradient Descent




Parameter Learning


- Solving a non-convex optimization.
- Iterative solution.
- Depends on the initialization.
- Convergence to a local optima.
- Judicious choice of learning rate



Least Squares versus Logistic Regression



Least Squares versus Logistic Regression




Bishop - PRML book (Chap 4)

Neural Networks


Perceptron Algorithm

Perceptron Model [McCulloch, 1943, Rosenblatt, 1957]




Targets are binary classes $[-1, 1]$

What if the data is not
linearly separable



Multi-layer Perceptron

Multi-layer Perceptron [Hopfield, 1982]




$$\mathbf{v}^2 = \psi \left(\mathbf{W}^2 \phi(\mathbf{W}^1 \mathbf{x} + \mathbf{b}^1) + \mathbf{b}^2 \right)$$

ϕ non-linear function (*tanh, sigmoid*)

ψ thresholding function

Neural Networks

Multi-layer Perceptron [Hopfield, 1982]




- Useful for classifying **non-linear data boundaries** - non-linear class separation can be realized given enough data.


Neural Networks

Types of Non-linearities ϕ


tanh



sigmoid



ReLu



Cost-Function

Mean Square Error

$$J_{MSE} = \sum_{i=1}^M ||\mathbf{v}_i - \mathbf{y}_i||^2$$

Cross Entropy

$$J_{CE} = - \sum_{i=1}^M \mathbf{y}_i^T \log(\mathbf{v}_i)$$

$\mathbf{y}_i$ are the desired outputs

Learning Posterior Probabilities with NNs

Choice of target function ψ

- Softmax function for classification

$$\psi(v_i) = \frac{e^{v_i}}{\sum_i e^{v_i}}$$

- Softmax produces positive values that sum to 1
- Allows the interpretation of outputs as posterior probabilities